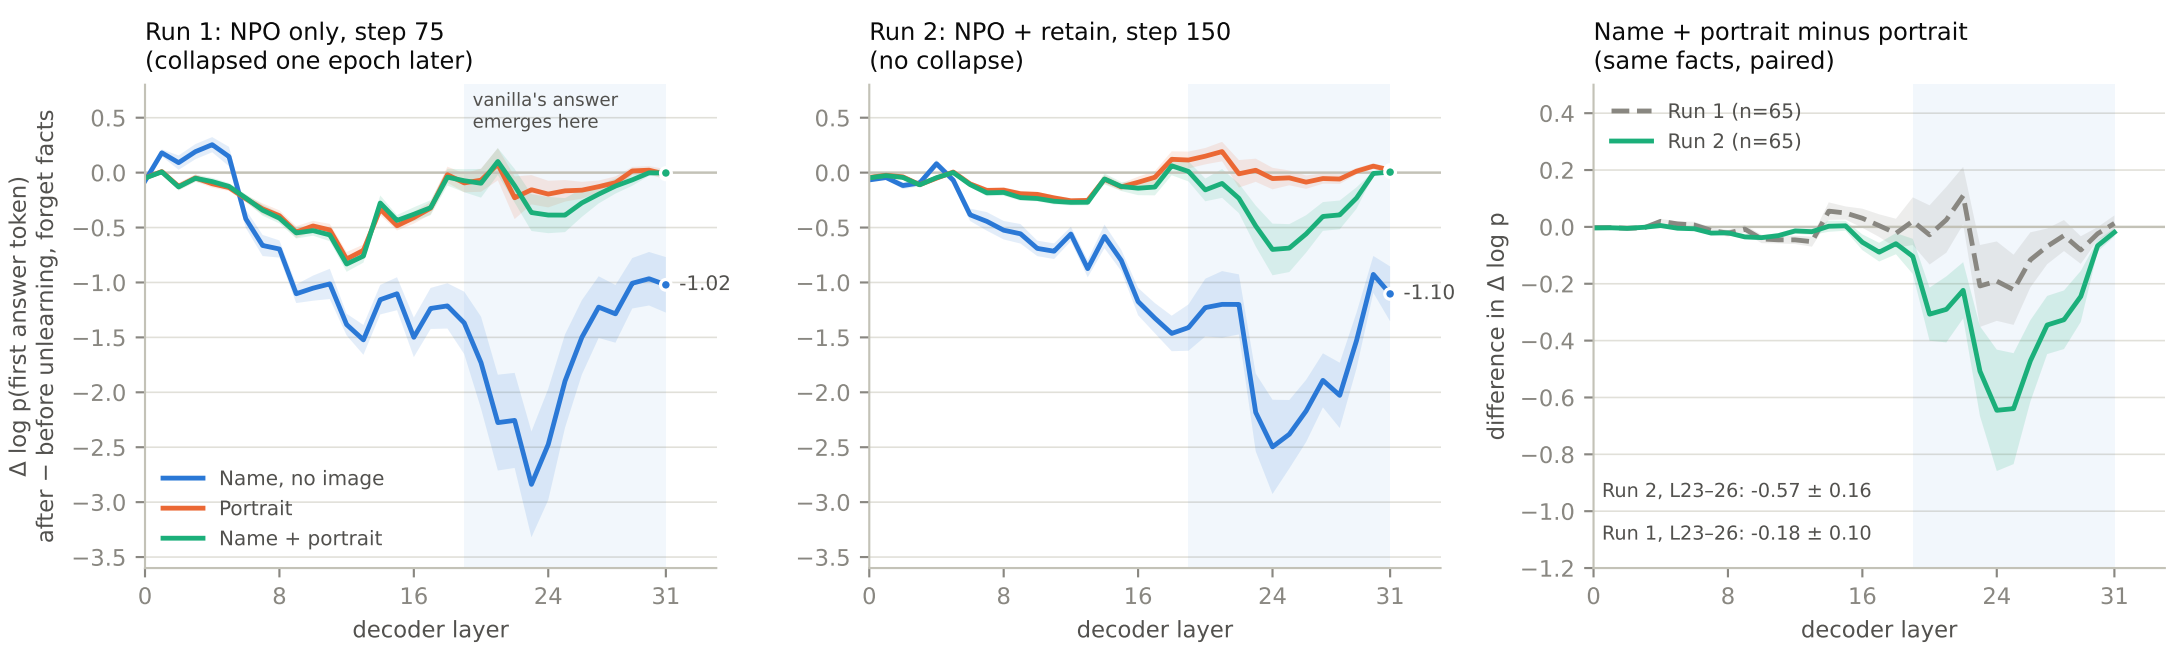
Run 1: NPO only, step 75
(collapsed one epoch later)
vanilla's answer
emerges here
Δ log p(first answer token)
after − before unlearning, forget facts
Name, no image
Portrait
Name + portrait
-1.02
decoder layer
Run 2: NPO + retain, step 150
(no collapse)
-1.10
decoder layer
Name + portrait minus portrait
(same facts, paired)
Run 1 (n=65)
Run 2 (n=65)
difference in Δ log p
Run 2, L23–26: -0.57 ± 0.16
Run 1, L23–26: -0.18 ± 0.10
decoder layer